

Shubham Reddy

AI Engineer | LangGraph · RAG · LangChain · Gemini API | TCS | Open to AI/Automation Roles

Pune, India 📍 +91-9172587538 ✉ shubhamreddy9172@gmail.com

LinkedIn · GitHub · Portfolio

Summary

AI Engineer with 2 years of production engineering experience at TCS. Built and deployed an AI agent in a live enterprise environment that auto-generates incident closure comments from application logs, reducing manual documentation time significantly. Developed a RAG-based onboarding system that reduced employee time-to-productivity by 85%. Experienced in LLM integration, prompt engineering, and end-to-end AI application development using LangChain, Gemini API, and NVIDIA NIM. Seeking AI engineering roles where I can build production-grade AI systems.

Technical Skills

Core: LangGraph, LangChain, RAG Systems, Gemini API, Python

Supporting: ChromaDB, Docker, SQL, GCP, Jenkins, LangSmith

Exposure: LlamaIndex, LangSmith, MCP, A2A Protocol, NVIDIA NIM

No-Code: Microsoft Copilot Studio

Experience

System Engineer

May 2024 – Present

Tata Consultancy Services (TCS) – Pune

- Reduced actionable alert noise by ~40%, saving the support team 2 hrs/day in manual triage.
- Executed 50+ SQL data updates across production databases to extract and validate data, supporting incident analysis and application health reporting.
- Maintained Jenkins pipelines across 3 production environments, ensuring zero-downtime deployments for enterprise applications.
- Monitored 5+ GCP services with 99.9% uptime tracking, and owned daily health reporting across the application stack.
- Deployed and validated containerized application services using Docker in controlled pre-production environments.

Projects

Digital Control Room (AI Portfolio) | github.com/shubhamr9172/sr-portfolio

Next.js, FastAPI, NVIDIA NIM, Framer Motion

- Developed a high-fidelity "mission control" portfolio featuring live interactive AI demos instead of static screenshots.
- Engineered an **Adaptive Interface** demo that enforces STAR-method constraints on user input via dynamic prompt engineering.
- Built an **L2 Diagnostic Terminal** using a sequential pipeline of Nemotron Content Safety and Nemotron-120B for log analysis and SOP generation.
- Implemented a **Logistics Reasoning Engine** utilizing Chain-of-Thought (CoT) to extract and analyze unstructured invoice data with high precision.
- Designed a responsive **3D Neural Topography** hero section using Canvas API and Framer Motion, optimized for high-performance visual impact.

Enterprise Incident AI Agent

Microsoft Copilot Studio — TCS, 2025

- Built no-code AI agent in Microsoft Copilot Studio that auto-generates ITIL-compliant incident closure comments from application logs.
- Achieved active adoption with the agent used daily by an 8-member support team for incident analysis.
- Reduced manual documentation time by ~60% by designing conversation flows and system prompts to enforce ITIL-compliant structures.

Zenith-LaTeX | <https://github.com/shubhamr9172/Zenith-LaTeX> | *React, Gemini API, NVIDIA Gemma 2, LaTeX*

- Architected a 100% serverless React application that transforms raw LaTeX code into ATS-compliant

PDFs, reducing the manual resume update cycle by over 70%.

- Optimized operational costs to \$0/month by bypassing CORS restrictions via a client-side hidden form submission pattern to the TeXLive API.
- Integrated Gemini & NVIDIA Gemma 2 APIs with a custom system instruction layer that enforces the STAR method and Action Verbs, ensuring high-impact output.
- Resolved 100% of editor crashes by engineering custom model lifecycle management within the Monaco Editor integration.

Corporate Onboarding Assistant (AI-RAG) | github.com/shubhamr9172/corporate-onboarding-assistant

Tech stack: LangChain, ChromaDB, Gemini API, Python

- Built RAG-based corporate onboarding assistant using LangChain, ChromaDB, and Gemini API that reduced the employee onboarding lifecycle from 1 week to 1 day (85% reduction in time-to-productivity).
- Built a high-performance retrieval pipeline, automating access request guidance and reducing lead support queries by 40%.
- Developed a role-aware personalization engine to dynamically filter knowledge retrieval based on user profile, tech stack, and project requirements.
- Designed a scalable document ingestion engine for multi-format enterprise policies, leveraging ChromaDB for high-accuracy semantic search.

Certifications

Microsoft | Applied Skill: Create Agents in Microsoft Copilot Studio (Mar 2026)

Google | Introduction to Generative AI

Intel | AI Aware Program

Education

Bachelor of Engineering – Computer Science 2023

Marathwada Mitra Mandal's Institute of Technology, Pune

Diploma – Information Technology 2020

Puranmal Lahoti Government Polytechnic, Latur

Additional

Created educational content on Generative AI and Prompt Engineering; authored technical blogs on AI, automation, and software engineering.